Jupyter Notebook Assignment Report

Student: Nishal Sukumar
Assignment: $Nishal_Sukumar_Assignment_1$
Date: 2025-04-11 19:29

Cell 1 (code)

```
\textbackslash{}# Import libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import fetch\textbackslash{}_california\textbackslash{}
    _housing
from sklearn.model\textbackslash{}_selection import train\textbackslash{}_test\
    textbackslash{}_split
from sklearn.linear\textbackslash{}_model import LinearRegression
from sklearn.metrics import mean\textbackslash{}_squared\textbackslash{}_error, r2
    \textbackslash{}_score
```

Cell 2 (code)

```
\textbackslash{}# Load California housing dataset
california = fetch\textbackslash{}_california\textbackslash{}_housing()
df = pd.DataFrame(california.data, columns=california.feature\textbackslash{}
    _names)
df['Target'] = california.target
```

Cell 3 (markdown)

\ —- EDA —-

Cell 4 (code)

```
\textbackslash{}# Basic Info
print("Shape of dataset:", df.shape)
print("\textbackslash{}nFirst 5 rows:\textbackslash{}n", df.head())
print("\textbackslash{}nSummary statistics:\textbackslash{}n", df.describe())
```

```
stdout: Shape of dataset: (20640, 9)

First 5 rows:
    MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  \
0  8.3252      41.0  6.984127   1.023810       322.0  2.555556     37.88
1  8.3014      21.0  6.238137   0.971880      2401.0  2.109842     37.86
2  7.2574      52.0  8.288136   1.073446       496.0  2.802260     37.85
```

```
3  5.6431      52.0  5.817352   1.073059       558.0  2.547945     37.85
4  3.8462      52.0  6.281853   1.081081       565.0  2.181467     37.85

   Longitude  Target
0    -122.23   4.526
1    -122.22   3.585
2    -122.24   3.521
3    -122.25   3.413
4    -122.25   3.422

Summary statistics:
             MedInc      HouseAge      AveRooms     AveBedrms    Population  \
count  20640.000000  20640.000000  20640.000000  20640.000000  20640.000000
mean       3.870671     28.639486      5.429000      1.096675   1425.476744
std        1.899822     12.585558      2.474173      0.473911   1132.462122
min        0.499900      1.000000      0.846154      0.333333      3.000000
25%        2.563400     18.000000      4.440716      1.006079    787.000000
50%        3.534800     29.000000      5.229129      1.048780   1166.000000
75%        4.743250     37.000000      6.052381      1.099526   1725.000000
max       15.000100     52.000000    141.909091     34.066667  35682.000000

            AveOccup      Latitude     Longitude        Target
count  20640.000000  20640.000000  20640.000000  20640.000000
mean       3.070655     35.631861   -119.569704      2.068558
std       10.386050      2.135952      2.003532      1.153956
min        0.692308     32.540000   -124.350000      0.149990
25%        2.429741     33.930000   -121.800000      1.196000
50%        2.818116     34.260000   -118.490000      1.797000
75%        3.282261     37.710000   -118.010000      2.647250
max     1243.333333     41.950000   -114.310000      5.000010
```

Cell 5 (code)

```
\textbackslash{}# Check for missing values
print("\textbackslash{}nMissing values:\textbackslash{}n", df.isnull().sum())
```

```
stdout:
Missing values:
 MedInc        0
HouseAge       0
AveRooms       0
AveBedrms      0
Population     0
AveOccup       0
Latitude       0
Longitude      0
Target         0
dtype: int64
```

Cell 6 (code)

```
\textbackslash{}# Correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```

```
<Figure size 1000x800 with 2 Axes>
```

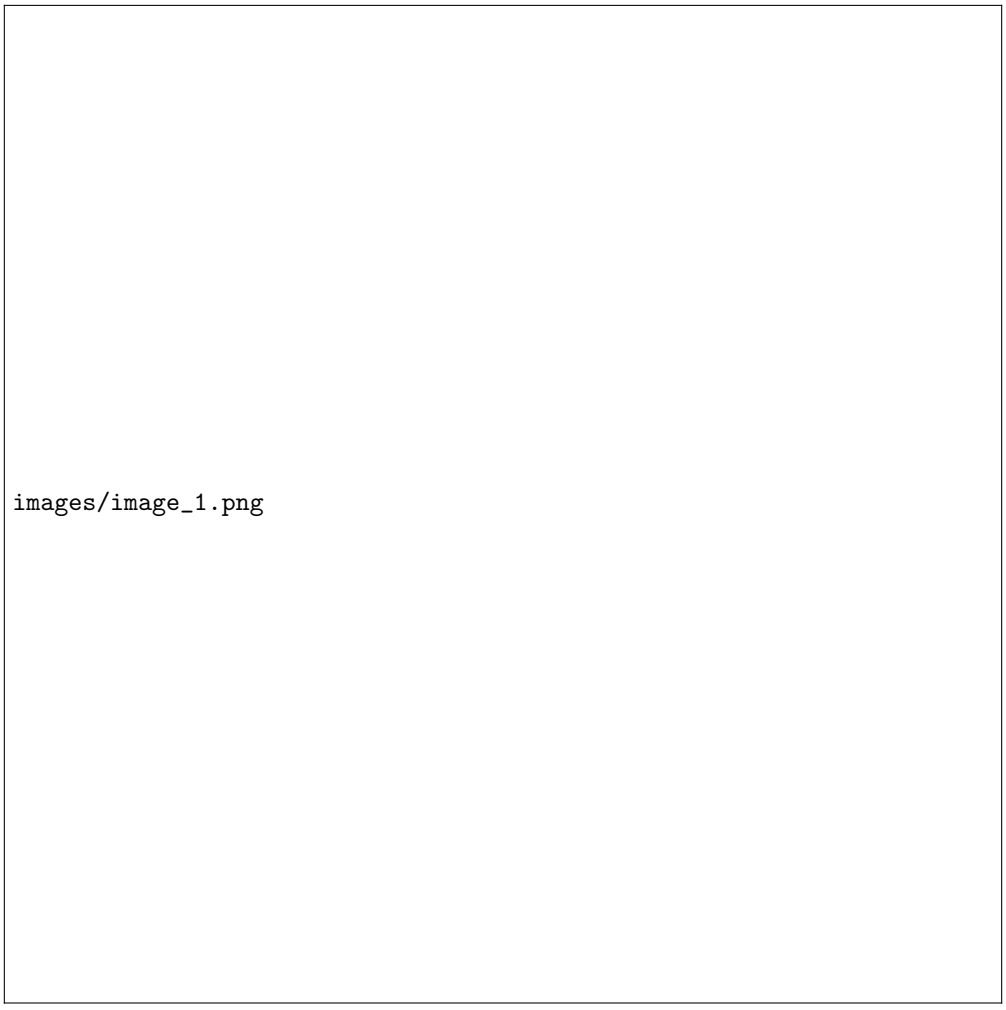

Figure 1: Output Image 1

Cell 7 (code)

```
\textbackslash{}# Distribution of the target variable
sns.histplot(df['Target'], kde=True)
plt.title("Distribution of Target (Median House Value)")
plt.show()
```

```
<Figure size 640x480 with 1 Axes>
```

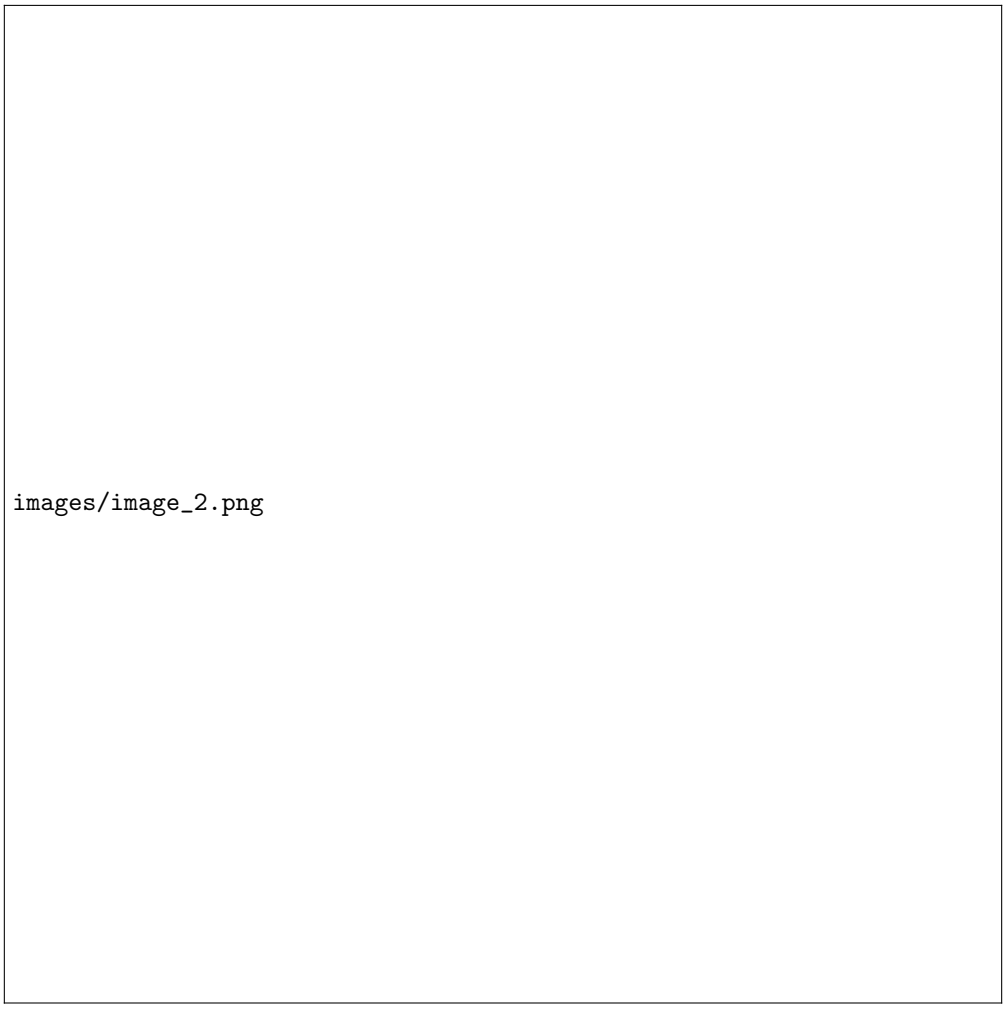

Figure 2: Output Image 2